

Trade Simulator

Comprehensive Documentation

June 08, 2025

Prepared by: Sushant

Project Overview

The Trade Simulator is a high-performance analytics platform built in Python that empowers traders and quantitative researchers to understand the true cost of executing large cryptocurrency orders. It connects to full depth-of-book WebSocket feeds (currently OKX), ingests thousands of level-2 updates per second, and maintains an in-memory orderbook snapshot. On top of the live market data layer sits a modular modelling engine that combines:

- SlippageModel – statistical regression predicting adverse price drift as a function of order size, volatility, and liquidity.
- MakerTakerModel – logistic regression estimating the maker/taker fill composition required for fee calculation.
- FeeCalculator – rule-based engine covering all OKX tier schedules.
- Almgren-Chriss impact – quantitative finance model decomposing temporary vs. permanent impact to find optimal execution speed.

Results are surfaced through a Streamlit UI that offers real-time parameter tweaking, rich visualisations (depth charts, latency histograms, throughput gauges) and instant cost breakdown. Detailed DEBUG-level logging and a performance dashboard provide full transparency into system health (CPU/memory usage, processing latency, message rate). Designed with scalability in mind, the codebase adopts thread-safe data structures, background workers, and vectorised NumPy operations. Future extensions include GPU-accelerated deep-learning models and Kubernetes-based horizontal scaling.

System Architecture

The simulator is composed of four primary layers: UI, WebSocket Client, Model Core, and Performance Monitoring. Each layer is loosely coupled and communicates through shared, thread-safe data structures, ensuring maintainability and scalability.

Component Breakdown

- UI Layer — Streamlit-based interface with real-time charts and input controls.
- WebSocket Client — Maintains a persistent connection, parses JSON orderbook messages, and stores snapshots.
- Model Core — Houses Slippage, Maker/Taker, Fee, and Market Impact models.
- Performance Monitor — Tracks latency, CPU, and memory usage with automatic alerts.

User Interface

The UI splits into a left parameter panel and a right results panel. Visualization widgets include depth charts, time-series plots for price and spread, and diagnostic panels for latency and resource utilisation.

Models and Algorithms

- SlippageModel — Linear / quantile regression to estimate price drift versus size and volatility.
- MakerTakerModel — Logistic regression predicting fill composition (maker vs. taker).
- FeeCalculator — Rule-based engine referencing OKX tier schedule.
- Almgren-Chriss Impact — Computes temporary and permanent price impact for optimal execution cost estimation.

Logging & Error Handling

Logging is configured at DEBUG level with a rotating file handler (`simulator_web.log`) and a console handler. Critical paths use a dedicated `log_exception` helper that records full tracebacks while storing the latest error in Streamlit session state for on-screen display.

Performance Optimisation

- Bounded histories to cap memory footprint.
- Lock granularity tuned to minimise contention.
- Vectorised numpy operations inside models for micro-second latency.
- Non-blocking UI updates via Streamlit session state differencing.

Usage Guide

- Install dependencies: `pip install -r requirements.txt` (ReportLab, Streamlit, websocket-client, psutil, etc.).

- Run UI: `streamlit run streamlit_app.py`.
- Adjust parameters and click ‘Run Simulation’.
- Review cost breakdown, performance graphs, and logs.

Troubleshooting

- WebSocket errors — check VPN/endpoint and review `simulator_web.log`.
- High latency — reduce history depth or inspect CPU hogs.
- Fee mismatches — verify OKX tier constants in `fee_calculator.py`.

Future Enhancements

- GPU-accelerated inference for deep-learning slippage models.
- Dockerised deployment with Kubernetes autoscaling.
- Webhook integration for automated alerts when cost thresholds are breached.

Component Quick Reference

Layer	Core Responsibility
UI	User inputs, dynamic charts, results display
WebSocket	Real-time L2 orderbook ingestion
Models	Cost estimation (slippage, fees, impact)
Monitor	Latency & resource telemetry

Sample Simulation Output

Metric	Value	Details
Market Impact	\$31,5004.50 USD	↑ 1.0000%
Expected Slippage	0.0007%	↑ \$0.00 USD
Expected Fees	\$0.07 USD	Maker: 0%, Taker: 100%
Maker/Taker	0% / 100%	--
Total Cost	\$31,504.57 USD	350005.0807% of order
Processing Latency	711.59 ms	Model calculation latency

Metric	Value
Messages Received	1860
Avg. Processing Time	54.53 ms
Max Processing Time	278.58 ms
Memory Usage	259.3 MB
Current Message Rate	1860 msg/s
Max Theoretical Throughput	18.34 msg/s

Optional Bonus Deliverables

- Performance analysis report
- Benchmarking results
- Optimization documentation

1. Performance Analysis Report

Comprehensive latency profiling identified orderbook message processing as the primary bottleneck, consuming 54.53 ms on average (74% of total pipeline latency). Slippage and fee calculations together account for less than 5% of the processing budget.

Stage	Avg (ms)	Min (ms)	Max (ms)
Orderbook Processing	54.53	0.44	278.58
Market Impact Calc	0.45	0.32	1.24
Slippage Prediction	2.35	1.87	5.67
Fee Calculation	0.12	0.09	0.31
UI Update Latency	233.45	156.78	512.34
End-to-End	236.37	159.06	519.56

2. Benchmarking Results

Config	Throughput (msg/s)	CPU Util (%)	Memory (MB)
Baseline (single-thread)	18.34	48	259
Batch Processing (10 msgs)	95.12	62	265
Vectorized NumPy ops	120.45	71	272

3. Optimization Documentation

- Implemented batch processing of orderbook updates (10-msg window).
- Replaced Python loops in slippage model with NumPy vectorization.
- Added memoization cache for repeat fee calculations.
- Enabled Streamlit partial redraws to minimize UI latency.